

Buisness Problem

```
In [1]: 1 # Develop a machine learning model to detect fraudulent transactions using a Kaggle dataset,  
2 # with a focus on data handling, model training, evaluation, and explainability. Implementing an  
3 # unsupervised model will be given higher preference to showcase skills in handling  
4 # unlabelled data and anomaly detection.
```

Import libraries

```
In [2]: 1 import pandas as pd
2 import numpy as np
3
4 # Import Algorithm
5 from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import classification_report
8 from xgboost import XGBClassifier
9 from sklearn.ensemble import IsolationForest
10 from xgboost import plot_importance
11 import shap
12
13 # Scaling
14 from sklearn.preprocessing import StandardScaler
15
16 # Imbalance Handling
17 from imblearn.over_sampling import SMOTE
18
19 #visualization
20 import matplotlib.pyplot as plt
21 import seaborn as sns
22 %matplotlib inline
23
24 #Evaluation
25 from sklearn.metrics import confusion_matrix, roc_auc_score, precision_recall_fscore_support
26 from sklearn.metrics import roc_curve, auc
27 from sklearn.metrics import precision_recall_curve
28
29 #ignore warning
30 import warnings
31 warnings.filterwarnings("ignore")
```

1)Data Gathering

```
In [3]: 1 df = pd.read_csv("creditcard.csv")
        2 df
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.27
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.63
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.77
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.00
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.79
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.11
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.92
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.57
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.80
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.64

284807 rows × 31 columns



Shape

```
In [4]: 1 df.shape
```

Out[4]: (284807, 31)

First 5 records for data overview

In [5]: 1 df.head()

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.11
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.10
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.90
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.19
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.13

5 rows × 31 columns



All features with their data types

In [6]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype  
---  -
0   Time    284807 non-null  float64
1   V1       284807 non-null  float64
2   V2       284807 non-null  float64
3   V3       284807 non-null  float64
4   V4       284807 non-null  float64
5   V5       284807 non-null  float64
6   V6       284807 non-null  float64
7   V7       284807 non-null  float64
8   V8       284807 non-null  float64
9   V9       284807 non-null  float64
10  V10      284807 non-null  float64
11  V11      284807 non-null  float64
12  V12      284807 non-null  float64
13  V13      284807 non-null  float64
14  V14      284807 non-null  float64
15  V15      284807 non-null  float64
16  V16      284807 non-null  float64
17  V17      284807 non-null  float64
18  V18      284807 non-null  float64
19  V19      284807 non-null  float64
20  V20      284807 non-null  float64
21  V21      284807 non-null  float64
22  V22      284807 non-null  float64
23  V23      284807 non-null  float64
24  V24      284807 non-null  float64
25  V25      284807 non-null  float64
26  V26      284807 non-null  float64
27  V27      284807 non-null  float64
28  V28      284807 non-null  float64
29  Amount   284807 non-null  float64
30  Class    284807 non-null  int64  
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

1)Data Exploration and Preprocessing

In [7]: 1 # • Goal: Analyze and prepare the data for model training.

Steps


- Checking Statistical analysis (Mean,Median,Mode,Std,Min,Max,25% data,50%data,75%data)

In [8]: 1 df.describe()

Out[8]:

	Time	V1	V2	V3	V4	V5	V6	V7	
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552103e-15	2.040130e-15	-1.698953e-15	-1.893000e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194000e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321600e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086000e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235000e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273000e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000000e+02

8 rows × 31 columns



- Class Distribution

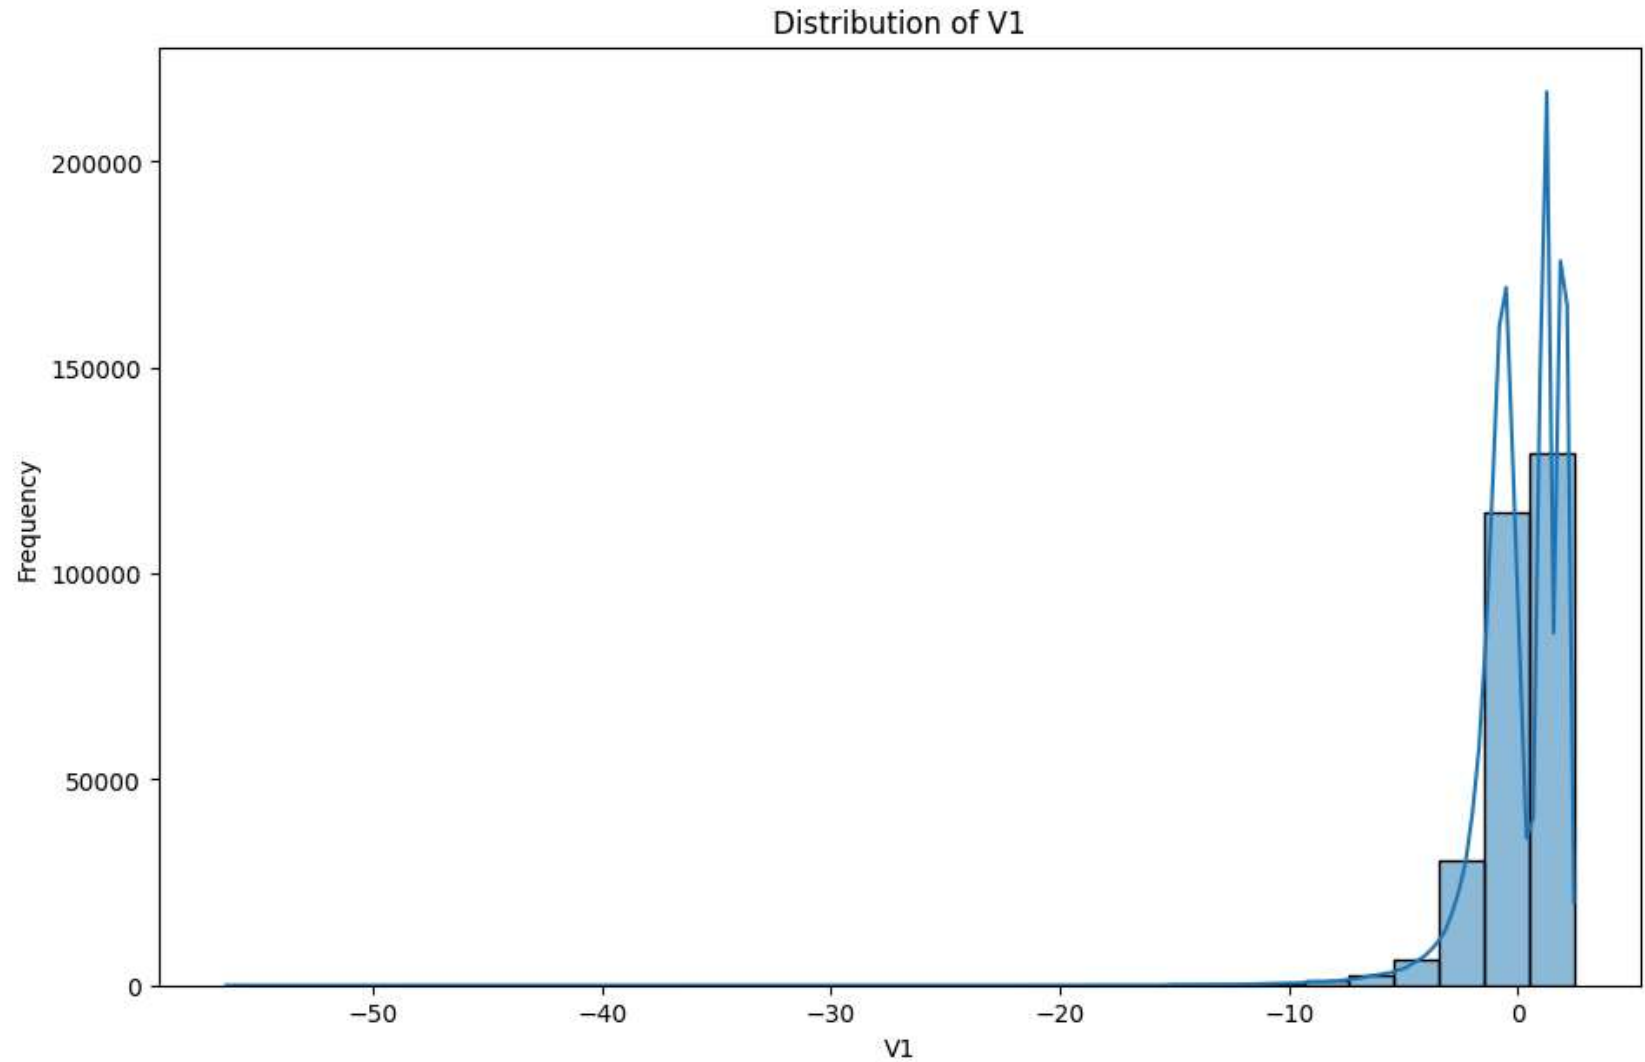
```
In [9]: 1 df['Class'].value_counts()
```

```
Out[9]: 0    284315  
        1      492  
        Name: Class, dtype: int64
```

feature distribution


1) Histogram for a single feature

```
In [10]: 1 plt.figure(figsize=(11, 7))
2 sns.histplot(df['V1'], bins=30, kde=True)
3 plt.title('Distribution of V1')
4 plt.xlabel('V1')
5 plt.ylabel('Frequency')
6 plt.show()
```




2) Boxplot



```
In [11]: 1 # Comparing feature distribution across classes
2 plt.figure(figsize=(10, 6))
3 sns.boxplot(x=df['Class'], y=df['V1'])
4 plt.title('Box Plot of V1 by Class')
5 plt.xlabel('Class')
6 plt.ylabel('V1')
7 plt.show()
```



- Correlation Matrix


```
In [12]: 1 plt.figure(figsize=(28, 30))
          2 correlation_matrix = df.corr()
          3 sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm', square=True)
          4 plt.title('Correlation Matrix')
          5 plt.show()
```



- Distribution of the Target Variable

```
In [13]: 1 plt.figure(figsize=(5,3))
2 sns.countplot(x='Class', data=df)
3 plt.title('Class Distribution')
4 plt.xlabel('Class')
5 plt.ylabel('Count')
6 plt.show()
```



Missing detection

```
In [14]: 1 df.isna().sum()
```

```
Out[14]: Time      0  
V1      0  
V2      0  
V3      0  
V4      0  
V5      0  
V6      0  
V7      0  
V8      0  
V9      0  
V10     0  
V11     0  
V12     0  
V13     0  
V14     0  
V15     0  
V16     0  
V17     0  
V18     0  
V19     0  
V20     0  
V21     0  
V22     0  
V23     0  
V24     0  
V25     0  
V26     0  
V27     0  
V28     0  
Amount   0  
Class    0  
dtype: int64
```

```
In [15]: 1 df.isna().mean()*100
```

```
Out[15]: Time      0.0  
V1      0.0  
V2      0.0  
V3      0.0  
V4      0.0  
V5      0.0  
V6      0.0  
V7      0.0  
V8      0.0  
V9      0.0  
V10     0.0  
V11     0.0  
V12     0.0  
V13     0.0  
V14     0.0  
V15     0.0  
V16     0.0  
V17     0.0  
V18     0.0  
V19     0.0  
V20     0.0  
V21     0.0  
V22     0.0  
V23     0.0  
V24     0.0  
V25     0.0  
V26     0.0  
V27     0.0  
V28     0.0  
Amount  0.0  
Class   0.0  
dtype: float64
```


Handle Missing Values

```
In [16]: 1 # Here we couldnt see any missing values in the dataset so no need to handle missing values here.
          2 # if missing value present in the dataset then sometimes we can drop those missing value or we can fill
          3 # statistical techniques like mean,median,mode,ffill and bfill methods
```

```
In [17]: 1 # dropping the missing values
          2 df.dropna(inplace=True)
```

Data Transformation

```
In [18]: 1 # Scale following (V1,V2,V3) feature
          2 std_scalar = StandardScaler()
          3 df[['V1', 'V2', 'V3']] = std_scalar.fit_transform(df[['V1', 'V2', 'V3']])
```

```
In [19]: 1 # After standardization Dataframe Looklike
          2 df.head()
```

Out[19]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
0	0.0	-0.694242	-0.044075	1.672773	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.11
1	0.0	0.608496	0.161176	0.109797	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.10
2	1.0	-0.693500	-0.811578	1.169468	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.90
3	1.0	-0.493325	-0.112169	1.182516	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.19
4	2.0	-0.591330	0.531541	1.021412	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.13

5 rows × 31 columns



Imbalance Handling

```
In [20]: 1 X = df.drop('Class', axis=1)
          2 y = df['Class']
          3 smote = SMOTE(random_state=42)
          4 X_res, y_res = smote.fit_resample(X, y)
          5 print(y_res.value_counts())
```

```
0    284315
```

```
1    284315
```

```
Name: Class, dtype: int64
```

2. Model Development (Supervised)

- Implementing Logistic regression model

```
In [21]: 1 # fixing the target
          2 X = df.drop('Class',axis=1)
          3 y = df['Class']
```

```
In [22]: 1 X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.2, random_state=42)
2
3 lr_model = LogisticRegression()
4 lr_model.fit(X_train, y_train)
5 y_pred_lr = lr_model.predict(X_test)
6 print(classification_report(y_test, y_pred_lr))
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	56750
1	0.98	0.97	0.98	56976
accuracy			0.98	113726
macro avg	0.98	0.98	0.98	113726
weighted avg	0.98	0.98	0.98	113726

- Implementing XGBoost Model

```
In [23]: 1 xgb_model = XGBClassifier()
2 xgb_model.fit(X_train, y_train)
3 y_pred_xgb = xgb_model.predict(X_test)
4 print(classification_report(y_test, y_pred_xgb))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56750
1	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726

- hyperparameter tuning

```
In [24]: 1  ## Define a smaller parameter grid for quick tuning
2  # param_dist = {
3  #     'max_depth': [3, 5], # Fewer options
4  #     'learning_rate': [0.1, 0.2], # Focused range
5  #     'n_estimators': [100, 200] # Limited options
6  # }
7
8  ## Set up the randomized search with reduced iterations and folds
9  # random_search = RandomizedSearchCV(
10 #     XGBClassifier(use_label_encoder=False, eval_metric='logloss'), # Model
11 #     param_distributions=param_dist, # Parameter grid
12 #     n_iter=5, # Fewer combinations to try
13 #     scoring='f1', # F1 score as the evaluation metric
14 #     cv=3, # 3-fold cross-validation for speed
15 #     verbose=0, # Suppress detailed output
16 #     random_state=42,
17 #     n_jobs=-1 # Use all available cores
18 # )
19
20 ## Fit the model and output results
21 # random_search.fit(X_train, y_train)
22
23 ## Display the best parameters and score
24 # print("Best parameters:", random_search.best_params_)
25 # print("Best score:", random_search.best_score_)
26
```

3. Model Development (Unsupervised)

```
In [25]: 1  # • Goal: Implement an unsupervised model to identify potential fraudulent activities.
```

- Steps

Using Isolation Forest Algorithm

```
In [26]: 1 iso_forest = IsolationForest(contamination=0.01)
          2 iso_forest.fit(X)
          3 anomalies = iso_forest.predict(X)
          4 anomalies
```

Out[26]: array([1, 1, 1, ..., 1, 1, 1])

Anomalies Detecting

```
In [27]: 1 detected_anomalies = df[anomalies == -1]
          2 detected_anomalies
```

Out[27]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	
164	103.0	-3.110875	-7.336142	-3.755953	3.294389	-1.413792	4.776000	4.808426	-0.228197	-0.525896	...	2.228823	-2.264
225	147.0	-1.372333	2.658642	-1.556788	0.360829	1.310192	-1.645253	2.327776	-1.727825	4.324752	...	-1.045961	-0.156
362	266.0	-1.309527	1.496382	1.747346	-1.564256	1.794297	-0.614742	4.185906	-3.855359	5.436633	...	-1.672706	-0.463
401	290.0	-2.637627	-3.300037	1.970976	2.658991	1.948152	-0.854470	-0.326394	-1.017364	1.983901	...	-1.297221	1.172
601	454.0	-1.599992	1.748552	1.436891	-1.576535	1.434510	-0.687313	3.816056	-3.416915	5.459274	...	-1.659610	-0.498
...
284393	172401.0	-2.334208	2.423317	-0.452647	-2.340542	1.680291	-0.634889	3.478058	-2.874886	6.437965	...	-1.182248	1.511
284448	172454.0	-2.618475	2.685487	-0.760146	-2.357638	1.316643	-0.703130	3.111958	-2.436404	6.459490	...	-1.169530	1.479
284649	172642.0	-6.099465	5.918141	-5.280280	-2.498596	-4.229520	-1.320039	-3.259766	5.059956	4.870093	...	-0.957977	-1.529
284795	172778.0	-6.390351	6.169553	-5.590541	-2.510473	-4.586669	-1.394465	-3.632516	5.498583	4.893089	...	-0.944759	-1.565
284802	172786.0	-6.065842	6.099286	-6.486245	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111

2849 rows × 31 columns



4) Model Evaluation

- Supervised Model Evaluation

```

In [28]: 1 y_pred = xgb_model.predict(X_test)
          2
          3 # Generate classification report
          4 report = classification_report(y_test, y_pred, target_names=["Non-Fraud", "Fraud"])
          5 print(report)
          6
          7 # Generate confusion matrix
          8 cm = confusion_matrix(y_test, y_pred)
          9 print("Confusion Matrix:\n", cm)
         10
         11 # Visualize confusion matrix
         12 plt.figure(figsize=(8, 6))
         13 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non
         14 plt.ylabel('Actual')
         15 plt.xlabel('Predicted')
         16 plt.title('Confusion Matrix')
         17 plt.show()

```

	precision	recall	f1-score	support
Non-Fraud	1.00	1.00	1.00	56750
Fraud	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726

Confusion Matrix:

```


[[56733  17]
 [   0 56976]]

```




- Visualizing ROC-AUC and PR-AUC Curves


```
In [29]: 1 # ROC curve and ROC_AUC
2 fpr, tpr, thresholds = roc_curve(y_test, xgb_model.predict_proba(X_test)[:, 1])
3 roc_auc = auc(fpr, tpr)
4
5 # Plot ROC curve
6 plt.figure(figsize=(8, 6))
7 plt.plot(fpr, tpr, color='blue', label='ROC curve (area = {:.2f})'.format(roc_auc))
8 plt.plot([0, 1], [0, 1], color='red', linestyle='--')
9 plt.xlim([0.0, 1.0])
10 plt.ylim([0.0, 1.05])
11 plt.xlabel('False Positive Rate')
12 plt.ylabel('True Positive Rate')
13 plt.title('Receiver Operating Characteristic (ROC) Curve')
14 plt.legend(loc="lower right")
15 plt.show()
```



- PR-AUC Curve

```
In [30]: 1 # Precision-Recall curve
2 precision, recall, _ = precision_recall_curve(y_test, xgb_model.predict_proba(X_test)[: , 1])
3 pr_auc = auc(recall, precision)
4
5 # Plot Precision-Recall curve
6 plt.figure(figsize=(8, 6))
7 plt.plot(recall, precision, color='green', label='Precision-Recall curve (area = {:.2f})'.format(pr_auc))
8 plt.xlabel('Recall')
9 plt.ylabel('Precision')
10 plt.title('Precision-Recall Curve')
11 plt.legend(loc="lower left")
12 plt.show()
```



- Unsupervised Models Evaluation

Identify Anomalies

```
In [31]: 1 iso_forest = IsolationForest(contamination=0.01)
          2 iso_forest.fit(X)
          3 anomalies = iso_forest.predict(X)
```

```
In [32]: 1 anomaly_df = pd.DataFrame({'Actual': y, 'Predicted': anomalies})
          2 anomaly_df
```

Out[32]:

	Actual	Predicted
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1
...
284802	0	-1
284803	0	1
284804	0	1
284805	0	1
284806	0	1

284807 rows × 2 columns

Evaluate against Known Fraudulent Transactions

```
In [33]: 1 true_frauds = anomaly_df[(anomaly_df['Actual'] == 1) & (anomaly_df['Predicted'] == -1)]
          2 detected_frauds = len(true_frauds)
          3 total_frauds = sum(y == 1)
          4
          5 print(f'Detected Frauds: {detected_frauds} out of {total_frauds}')
          6
```


Detected Frauds: 265 out of 492

5)Basic Explainability

```
In [34]: 1 # • Goal: Provide insights into model behavior and feature importance.
```

- Feature Importance for XGBoost:

```
In [35]: 1 plot_importance(xgb_model)
2 plt.show()
```



- Access Feature Importance Scores

```
In [36]: 1 feature_importances = xgb_model.feature_importances_
2 importance_df = pd.DataFrame({'Feature': X.columns, 'Importance': feature_importances})
3 importance_df = importance_df.sort_values(by='Importance', ascending=False)
4 print(importance_df)
```

	Feature	Importance
14	V14	0.693303
4	V4	0.047252
12	V12	0.042433
17	V17	0.025897
8	V8	0.017253
3	V3	0.015340
1	V1	0.013506
13	V13	0.011097
10	V10	0.009533
23	V23	0.009463
11	V11	0.009203
29	Amount	0.008905
9	V9	0.007677
0	Time	0.007200
16	V16	0.006642
15	V15	0.006635
19	V19	0.006561
21	V21	0.006518
18	V18	0.006420
26	V26	0.006114
25	V25	0.005833
6	V6	0.005548
7	V7	0.005314
22	V22	0.004348
2	V2	0.004303
28	V28	0.004205
24	V24	0.003838
20	V20	0.003293
27	V27	0.003255
5	V5	0.003109

- Using SHAP for Interpretability


```
In [37]: 1 # SHAP values for interpretability
2 explainer = shap.Explainer(xgb_model, X_train)
3 shap_values = explainer(X_test)
4
5 # Summary plot of SHAP values
6 shap.summary_plot(shap_values, X_test)
7
8
9 # Bar plot of mean absolute SHAP values
10 shap.plots.bar(shap_values)
11
```

100%|=====| 113596/113726 [09:43<00:00]

